

The Engine and the *Instrument*

On the architectural and strategic distinction between Kincaid IQ and Rebel IQ — and why the separation between them is the point.

Two names. One system. Kincaid IQ is the engine that reasons about fiduciary exposure at the level of the contract clause and the individual claim. Rebel IQ is the instrument that carries that reasoning to the people who hold the obligation — and turns it into a decision, a disclosure, and a record a participant can be defended by. They are not competitors. They are strata.

CLASSIFICATION	Internal architecture & positioning · Board distribution
ENGINE ENTITY	SiriusB AI Data Sciences — IP-holding parent
OPERATING BRAND	Healthcare Rebel Alliance — FRAME™ trademark holder
SUBJECT SYSTEMS	Kincaid IQ · Rebel IQ™ product family
OFFICE	Office of the Chief Architect
REGISTER	Sober, analytical, participant-first

The misreading this paper exists to retire.

The ecosystem operated by SiriusB AI Data Sciences and the Healthcare Rebel Alliance is frequently described — by people who ought to know better — as “two products that do roughly the same thing.” The description is wrong in a way that matters, because it obscures the single most consequential decision in the architecture: the deliberate separation of the intelligence from the instrument, and the choice to make the boundary between them a contract rather than a convenience.

Kincaid IQ is the engine. It is a continuous, multi-tenant fiduciary-intelligence substrate that computes ERISA exposure at the granularity of the contract clause and the dispensed claim, and exposes that computation to the world only through a versioned, authenticated API. Rebel IQ is the instrument: the consumer-facing product family that takes the engine's reasoning to the benefits committee, the plan sponsor, the broker, and the consultant, and renders it as a decision that can be made and a record that can be defended.

The distinction is not cosmetic. It determines who builds what, who is exposed to which liability, who buys which product and for which reason, and — most consequentially — how the architecture protects the plan participant whose name appears nowhere in either system but whose money is the entire reason either system exists. What follows is an account of that separation: what each layer is, where the seam between them runs, and why the seam is not a weakness to be engineered away but the load-bearing idea of the whole enterprise.

Two names, *one system*.

Why the ecosystem carries two product identities, and why the boundary between them was made a contract on purpose.

Every software enterprise of any ambition eventually meets a structural fact it did not plan for: the thing that makes the product valuable and the thing the customer actually touches are not the same thing, and the instinct to collapse them into one thing degrades both. The intelligence wants to be deep, slow-moving, rigorously versioned, and indifferent to who is looking at it. The instrument wants to be fast, opinionated, beautiful at the exact moment of use, and obsessively attentive to one particular human with one particular job. Forced to live in a single codebase, each compromises the other. The deep system accretes product hacks; the product accretes intelligence debt. Both become harder to defend in the one venue where defensibility is the entire game — a courtroom, seven years from now.

This ecosystem made the opposite choice, early and on purpose. It split the intelligence from the instrument and made the line between them a contract. Kincaid IQ does the hard reasoning. Rebel IQ delivers it. Between them sits a versioned, authenticated API, and that API is not an implementation detail to be optimized later. It is the architectural decision from which nearly everything else in this paper follows.

Two consequences fall out of that decision immediately, and the remainder of this paper is in large part an unpacking of them.

The first: **Kincaid IQ is consumed as a service, not embedded as a library.** The agent ecosystem, the verifier discipline, the audit ledger, the litigation-grade provenance chain, and the rest of the fiduciary-intelligence substrate live upstream and remain the engine's problem. Rebel IQ does not import them, does not vendor them, and does not re-implement them. It calls them. The contract between the two products is an HTTPS interface — authenticated, version-pinned, governed by an explicit service-level agreement with stated deprecation timelines — and that interface is the only surface across which intelligence crosses from one layer to the other.

The second: **Rebel IQ is therefore free to be a product rather than a platform.** Because it does not carry the weight of the reasoning engine, it can spend its entire engineering budget on the experience of the person who has to act — the committee chair preparing for a quarterly meeting, the broker building a pitch, the consultant assembling a deliverable. It does not re-solve contract extraction or risk scoring or dollar quantification, because those are solved, upstream, and paid for. Its scope is the application layer above the API, which is a more tractable surface and a faster one to ship.

• A NOTE ON THE NAMES

The naming follows the architecture. Kincaid IQ carries the name of the engine — a single system, a single substrate. Rebel IQ is the operating-brand expression, and it spans a family rather than a single product: a flagship committee platform and two engagement-scoped consulting wedges, all delivering the same underlying intelligence through different doors. The shared “IQ” suffix is not coincidence and not marketing symmetry; it is a true statement that the intelligence is common. The differing prefixes are an equally true statement that the layers are not.

The API is not the boundary between two teams. It is the boundary between two kinds of correctness — the correctness the engine owes a magistrate, and the correctness the instrument owes a committee chair on a Tuesday.

ON THE SEAM

Kincaid IQ — the fiduciary *intelligence engine*.

A continuous reasoning substrate, exposed only through an API, owned by the IP-holding parent. The durable asset of the enterprise.

Kincaid IQ is, at its core, a system that computes fiduciary state continuously. It does not produce a report on request and then go quiet. It is running every minute of every day, holding a current, decomposable account of where a self-funded plan stands against the obligations that ERISA, the Consolidated Appropriations Act, and the Department of Labor collectively recognize as actionable. The competitive category it occupies does not yet have a settled name; the durable noun underneath it is *fiduciary intelligence*, and the architectural commitment that distinguishes it from a conventional governance-risk-compliance product is that it replaces the checklist with continuous reasoning.

The reasoning is performed by an ecosystem of ten specialized agents. They are not chatbots bolted onto a binder; they are substrate agents, each producing a typed, evidence-bound output that downstream consumers — including Rebel IQ — can depend on. The ecosystem is worth naming in full, because the breadth of it is the point.

THE KINCAID IQ TEN-AGENT ECOSYSTEM

PCIA	PBM Contract Intelligence. Extracts fiduciarily material terms from pharmacy-benefit and vendor contracts at the level of the individual clause — spread-pricing definitions, rebate-aggregator carve-outs, affiliated-pharmacy steerage — and flags each against a benchmark corpus.
FRSA	Fiduciary Risk Scoring. Maintains a continuous, decomposable risk score per tenant across explicit statutory dimensions, each mapped to a defensible anchor rather than a “vibe.” This is the gauge the CFO looks at first and the artifact the regulator subpoenas if things go wrong.
FLQA	Financial Leakage Quantification. Translates detected anomalies and term deviations into defensible dollar estimates of recoverable or avoidable spend — the figure the expert witness will eventually cite.
RADA	Claims Anomaly Detection. An ensemble — statistical, classical ML, deep, rule-based, and a constrained LLM pattern matcher — operating over normalized claim streams, binding each anomaly to the specific contract term it violates.
CCDA	Compliance Change Detection. Surveils the federal and state regulatory perimeter and translates 17th-grade regulatory prose into “for self-insured plans above 5,000 lives, this means you must, by this date, do this.”
SEMA	SEC EDGAR Monitoring. Watches the full filing universe for disclosures that materially intersect a tenant’s vendors and counterparties, extracting typed signals from the ~3,000 filings that cross the wire each day.
HLSA	Litigation Surveillance. Tracks novel-theory litigation and clusters emerging legal theories so that a tenant receives lead time on the cases that go on to spawn follow-on suits.
ENGA	Narrative Generation. Turns the substrate’s reasoning into committee-grade prose, under the strict constraint that every assertion cite a specific underlying finding. Its measure is how often a narrative drives a human action within seventy-two hours.
BOIA	Broker Opportunity Intelligence. Serves the broker segment — and carries the system’s most important ethical guardrail, the fiduciary-alignment score, discussed at length in Section VI.
GAA	Governance & Operations. The meta-agent: cost telemetry, drift detection, performance, and the cross-agent arbitration that surfaces conflicts between agent outputs as a product feature rather than a hidden failure.

Three properties of this engine deserve emphasis, because they are exactly the properties Rebel IQ inherits without rebuilding. First, the **verifier discipline**: every extraction is independently checked by a verifier sub-agent, and any output below a calibrated confidence threshold is routed to a human before it propagates downstream. Second, the **audit ledger**: every analysis carries full provenance — the prompt family, the model version, the reasoning trace, the source span on the page — written from the first day, so that an output produced in 2026 can be defended in 2031 when opposing counsel subpoenas the methodology. Third, **substrate neutrality**: the engine has no opinion about who is consuming it. It emits typed objects. What a benefits committee, a broker, or a consultant does with those objects is a question for the layer above.

And the engine belongs, deliberately, to **SiriusB AI Data Sciences** — the IP-holding parent. This is not a trivial bookkeeping fact. The engine is the durable, defensible asset of the entire enterprise; placing it in the IP-holding entity, distinct from any operating brand, is the corporate expression of the same instinct that placed it behind an API: the thing that is hard to build and expensive to reproduce should be insulated, structurally, from the fortunes of any one product that happens to sell it.

Rebel IQ — the instrument at the *point of decision*.

A consumer-facing product family under the Healthcare Rebel Alliance brand, built where a human has to act and then defend the action.

Where Kincaid IQ reasons, Rebel IQ delivers — at the precise point where a human being has to make a fiduciary decision and then stand behind it. It is the consumer-facing layer, operated under the Healthcare Rebel Alliance brand, and it is decidedly not a single product. It is a family, and the differences within the family are as important to understand as the difference between the family and the engine beneath it.

- 1 **Rebel IQ™ Fiduciary Insights** — the flagship. A committee-facing ERISA compliance platform on a quarterly cadence, built on the FRAME™ six-agent MVP. This is the platform a benefits committee logs into before it meets: the running fiduciary record, the risk posture, the disclosure surfaces, the prose a chair can read into the minutes. It is recurring, it is a platform, and it is the home of the methodology described in Section V.
- 2 **Rebel IQ™ Opportunity Analysis** — a per-engagement consulting wedge, document-only, delivered through Claude Code. Not a standing platform but a deliverable: a scoped analytical artifact produced for a specific engagement and handed over.
- 3 **Rebel IQ™ Fee Intelligence** — likewise a per-engagement, document-only consulting wedge. The sharp edge of the commercial wedge: the comparative fee analysis that historically only the mega-consultancies could deploy, produced for one client, one engagement, one decision.

The distinction inside the family is a distinction of cadence and form. **Fiduciary Insights is a recurring platform** the institution lives in; the two wedges are **engagement-scoped documents** the institution receives. The platform is how the relationship persists; the wedges are how it begins. A consultant earns trust with a Fee Intelligence document and converts it into a standing Fiduciary Insights relationship — the wedge opens the door, the platform keeps it open.

• WHAT REBEL IQ BUILDS, AND WHAT IT DOES NOT

The defining discipline of the product layer is restraint about scope. Rebel IQ builds the application: the committee experience and the surfaces it presents; the book-level fiduciary memory that records decision rationale, the alternatives considered, and the prior-advisor context, so that institutional knowledge survives the turnover of the people who held it; the cross-tenant benchmark network whose value compounds as the corpus deepens; the compensation-disclosure surfaces mandated by the Consolidated Appropriations Act; and the running fiduciary record itself. What Rebel IQ does *not* build is the reasoning. It does not implement contract extraction, compensation scoring, or fee quantification, because those are Kincaid IQ outputs received over the API. The product layer turns the engine's intelligence into something a fiduciary can hold, read, sign, and — when the moment comes — defend.

And Rebel IQ belongs to the **Healthcare Rebel Alliance**, the operating brand — distinct, by design, from the IP-holding entity that owns the engine. The consumer-facing product is where the brand lives, the relationship is held, and the revenue flows; the engine is where the durable intelligence asset lives. Two layers, two homes.

The seam: dependency, not *duplication*.

A naive organization would have built the intelligence twice. This one built it once and consumes it cleanly — and names the risk that creates.

The single most important sentence in the relationship between the two systems is this: **Rebel IQ does not implement what Kincaid IQ implements**. The temptation to do otherwise is real and nearly universal — most organizations, handed a hard reasoning problem and a hard product problem, solve the reasoning problem twice, once badly inside the product. This ecosystem refused that path. The intelligence is built once, upstream, and consumed across a clean interface. The result is dependency, deliberately chosen, in place of duplication, accidentally accrued.

The engineering logic is straightforward and compounding. Consuming the engine rather than embedding it lets the product layer **ship faster**, because it inherits solved problems rather than re-solving them. It keeps the product's **codebase smaller**, because the verifier loops, the orchestration, and the audit machinery are not in it. And — least obvious but most valuable — it **isolates the product's litigation and audit surface** from the engine's far deeper one. When opposing counsel subpoenas methodology, the product's exposure is the thin application layer; the deep intelligence provenance is the engine's to defend, carried automatically through every analysis the product renders.

Honesty requires naming what the seam costs, because the discipline of this enterprise is to name failure modes rather than soften them. The dependency is real: if the engine's service-level agreement deteriorates, the committee's experience in the product degrades; if the API contract changes incompatibly, the product needs rework; if the engine pivots or is acquired, the supplier relationship enters uncertainty. None of this is denied. All of it is managed, and the management is itself part of the architecture.

THE DISCIPLINE OF THE SEAM

VERSION	A versioned API contract with explicit deprecation timelines, so a change upstream is a scheduled event the product layer can absorb — never a surprise that breaks it.
CACHE	Response caching at the Rebel IQ tier, so transient engine unavailability does not become committee-facing unavailability.
DEGRADE	Graceful degradation: when the engine is unreachable, the product falls to a read-only mode that preserves the audit ledger and every existing analysis. The record is never at the mercy of the link.
INDEPEND	Roadmap independence: neither product's schedule is permitted to hold the other hostage. The seam is a contract precisely so that each side can move without waiting on the other.

The lock that makes this work is procedural, and it is treated as a gating dependency rather than an optional milestone: the API contract is locked first, before the product is built on top of it, so that the foundation is settled before the structure rises. This is the inversion of the usual mistake, in which a product is built against an interface that is still in flux and then spends its life chasing the interface's changes.

Ten agents and three gates: substrate versus *methodology*.

The most common technical confusion — and the one this section exists to dissolve — is that Kincaid's ten agents and FRAME's six are the same kind of thing.

Because both Kincaid IQ and the FRAME™ methodology are described in terms of “agents,” the two are routinely conflated, and the conflation is the single most damaging misreading of the architecture after the engine-versus-instrument confusion itself. The ten agents and the six are different kinds of object. The ten are **substrate agents**: they produce typed intelligence — a contract term, a risk score, a dollar figure, a flagged anomaly. The six are **methodology agents**: they orchestrate. FRAME™ is a methodology layer, not a parallel or competing engine, and the distinction is the difference between muscle and choreography.

FRAME™ — the trademarked methodology held by the Healthcare Rebel Alliance — is the layer that turns the engine's raw reasoning into a structured fiduciary determination a committee can act on. Its minimum-viable form is six agents, and they correspond not to intelligence primitives but to the questions a fiduciary is actually obligated to answer.

THE FRAME™ SIX-AGENT MVP

- F-1** **Financing.** How the plan's money is financed and how it moves — the structure of the financial arrangement the fiduciary is responsible for understanding.
- F-2** **Service Provider Disclosure.** Whether the disclosures owed under ERISA §408(b)(2) — and broker compensation under CAA §408(b)(2)(B) — are in fact present, complete, and adequate.
- F-3** **Prohibited Transaction (three layers).** A three-layer test against ERISA §406 — party-in-interest screening, self-dealing detection, and affiliate-transaction analysis — the deepest of the six.
- F-4** **Cost Plus Unit Comparator.** Unit-level price comparison against a transparent Cost-Plus benchmark, the empirical heart of the “are these fees reasonable” question.
- F-5** **Fiduciary Gate.** The determination of whether the fiduciary *process* itself — not merely the outcome — meets the prudence standard.
- F-6** **Monitoring Infrastructure Assessment.** Whether ongoing monitoring infrastructure exists, because a prudent process is a continuing obligation, not a one-time event.

Above the agents sit **three gates**, which are the methodology's organizing logic: the **Conflict gate (G1)**, the **Reasonableness gate (G2)**, and the **Prohibited Transaction gate (G3)**. These map, respectively, to the conflict screen, to the reasonable-compensation standard of ERISA §408(b)(2), and to the prohibited-transaction prohibitions of §406. A finding does not simply accumulate in a pile; it passes through the gate that governs its statutory character, which is what allows the methodology to produce a determination rather than a data dump.

FRAME also imposes a discipline on how findings are **routed**, and the discipline is more subtle than it first appears. The methodology runs a *dual-stream output*: findings about *what a term says* — its content — route to their home substantive dimension, while findings about *whether a required term or disclosure exists at all* — its existence — route to a distinct Accountability dimension. The separation matters because the two failure modes are different in kind. A disclosed but unreasonable fee is a content problem; a disclosure that was simply never made is an accountability problem, and conflating the two would let a missing obligation hide inside a substantive analysis. For quantification, FRAME applies a conservative *captured-float convention* — a five-percent assumption on intermediary float, attributed pro-rata — so that the economic estimate is defensible by being deliberately understated rather than aggressive.

• HOW THE METHODOLOGY RIDES ON THE ENGINE

The relationship between the six and the ten is concrete, not abstract, and a single example makes it legible. Kincaid's **PCIA** extracts a service-provider compensation term from a contract; FRAME's **Service Provider Disclosure** agent then evaluates whether the disclosure obligation under §408(b)(2)(B) has been met, passing the result through the Reasonableness gate. Kincaid's **RADA** flags a claims anomaly consistent with affiliated-pharmacy steerage; FRAME's **Prohibited Transaction** layers then test it against the three prohibitions of §406, through gate G3. The engine supplies the evidence; the methodology supplies the verdict. Neither does the other's job, and the system is coherent precisely because they do not. The legal landscape that makes this choreography matter is not hypothetical — *Cunningham v. Cornell* (2025) reshaped how prohibited-transaction claims survive a motion to dismiss; *Tiara Yachts* (6th Cir. 2025) sharpened the stakes of plan-asset handling; and EBSA Field Assistance Bulletin 2026-01 moved the disclosure perimeter underfoot. The methodology exists to answer the questions these authorities actually ask.

The engine is the muscle; the methodology is the choreography. Ten agents reason about the world. Six agents decide what a fiduciary must do about it.

SUBSTRATE & METHODOLOGY

The dual-buyer architecture: where purpose *diverges*.

The two systems differ most sharply not in what they compute but in whom they serve — and the product layer resolves a genuine conflict of interest in the structure itself.

If the engine and the instrument differ in *what* they are, the place they differ most in *purpose* is the buyer. Kincaid IQ is substrate-neutral: it serves whoever holds the fiduciary obligation, with no opinion about their role. Rebel IQ is, by deliberate design, two-sided — and the two-sidedness is not a marketing posture but an engineered property of the system.

The product is sold into two postures at once. **Brokers buy it offensively:** a mid-market brokerage uses it to win accounts from incumbents whose advantage was built on analytical asymmetry, armed now with the comparative fee intelligence that historically only Mercer, Aon, and Willis Towers Watson could field at scale. **Plan sponsors buy it defensively:** the plan sponsor is the data subject, the holder of the obligation, the entity that must discharge a fiduciary duty and protect both itself and its participants. The same intelligence serves an offensive motion and a defensive one — and the obvious hazard is that a broker's interest is not always aligned with the client's.

This is where the architecture does something most platforms only promise. It resolves the conflict **structurally, in the system, rather than rhetorically, in a code of conduct**. The mechanism lives in the **BOIA** agent, which computes a **fiduciary-alignment score** on every broker-facing recommendation. Any recommendation that would benefit the broker at the client's expense — steering toward a higher-commission arrangement that scores worse on fiduciary risk, for instance — is *suppressed from the broker's interface entirely*, and surfaced instead to the underlying plan sponsor. The suppression clears a defined threshold and is logged to an audit the client can see; the broker is not merely discouraged from acting against the client's interest but is structurally prevented from even seeing the disaligned recommendation, unable to infer its existence through timing, count, or interface artifact.

Alignment is not a value statement in a sales deck. It is enforced at the data layer. A channel partner who asks for it to be softened has, in that request, identified themselves as the wrong partner.

THE FIDUCIARY-ALIGNMENT SPINE

The same logic governs the compensation-disclosure surface. Under CAA §408(b)(2)(B), a broker's compensation must be disclosed to the plan sponsor; the product makes that disclosure a *precondition of access* rather than an afterthought — a broker tenant declares all compensation arrangements as a condition of activating any grant against the client's data. The surface that renders this disclosure, and the audit ledger that records the broker's activity in the client's tenant, remain branded to the platform even under white-label, never to the broker. A broker who wishes to remove that mark is, again, telling the system something it should attend to.

This is the point at which the architecture stops being merely clever and becomes a statement of values rendered in software. It is structurally difficult, by design, to turn the tool against the participant. That difficulty is not an accident of implementation; it is a moral commitment encoded as architecture, and it is the spine of the entire broker product. The engine, neutral by nature, would serve a disaligned broker as readily as an aligned one. The instrument will not — and the reason the instrument can refuse is precisely that it is a separate layer, free to hold an opinion the substrate is not built to hold.

Cadence, surface, and *consumption*.

How the two systems feel different in the hand: one keeps the rhythm of the machine, the other the rhythm of the institution.

Architecture is felt, in the end, as cadence — the rhythm at which a system touches the people who use it. Here the two layers diverge cleanly, and the divergence is not incidental: each keeps the tempo appropriate to what it is.

Kincaid IQ keeps the rhythm of the machine. It is continuous; it computes fiduciary state every minute, never going quiet between events. It is consumed programmatically, through the API, and where its surfaces are exposed directly they take the form of an enterprise terminal — the **COMMAND, FIDUCIARY, CONTRACTS, INTELLIGENCE, ANOMALIES**, and **ARCHIVE** surfaces — in a register drawn from the instruments professionals already trust under pressure: deep navy and slate, gold reserved for the genuinely critical alert, a dark-mode terminal aesthetic that signals seriousness rather than consumer friendliness. It is a Bloomberg for fiduciary state, and it is always on.

Rebel IQ keeps the rhythm of the institution. The flagship, Fiduciary Insights, runs on a *quarterly* cadence — because the benefits committee meets quarterly, and the product's tempo is matched to the governance rhythm it serves rather than to the machine beneath it. It is a platform a committee logs into in the days before it convenes: the quarter's risk posture assembled, the findings routed through the gates, the narrative drafted into language a chair can read into the minutes. The two consulting wedges keep a different rhythm again — *per-engagement*, document-only, delivered through Claude Code as a scoped artifact for a particular client and a particular decision. There is no standing platform behind a wedge; there is a deliverable, produced when the engagement calls for it and handed over.

The pattern is clean enough to state in a sentence. The engine's cadence is the machine's — continuous, indifferent to the calendar. The instrument's cadence is the institution's — quarterly for the committee that lives in the platform, per-engagement for the consultant who arrives with a document. The engine never sleeps; the instrument speaks at the moment the fiduciary needs it to, and is quiet otherwise. A system that confused these — that tried to make the committee live inside a continuous terminal, or tried to compute fiduciary state only once a quarter — would serve neither the machine's logic nor the institution's. The separation lets each keep its proper time.

The corporate and IP *topology*.

The separation of layers is mirrored, deliberately, in a separation of entities and of intellectual property — so that the seam in the software is also a seam in the balance sheet.

The technical separation between engine and instrument is not left to live only in the codebase. It is mirrored, on purpose, in the corporate and intellectual-property structure, so that the boundary that protects each layer in software also protects each layer as an asset. The topology is worth laying out explicitly, because it is the legal expression of the same idea that runs through everything above.

ENTITY	ROLE IN THE SYSTEM	WHAT IT HOLDS
SIRIUSB AI DATA SCIENCES	Parent, IP-holding	Kincaid IQ — the engine. The durable, defensible intelligence asset, insulated from any single product's fortunes.
KINCAID RMC	Sister operating entity	The operating counterpart on the engine side of the house.
HEALTHCARE REBEL ALLIANCE	Operating brand	The Rebel IQ™ product family and the FRAME™ methodology trademark (Wisconsin registered, federal pending).
SYNOVIENT	Tokenization partner	A Phase-2 extension — CAA §724 plan-data title tokenization and FRAME™ IP-protection tokenization. Not in MVP scope.

Two features of this topology repay attention. First, the engine sits in the **IP-holding parent**, structurally distinct from the operating brand that sells the consumer product. This is the corporate analogue of placing the engine behind an API: the asset that is hardest to build and most expensive to reproduce is insulated from the operational and reputational volatility of any one product line. If the consumer brand were to pivot tomorrow, the engine would endure as an asset in its own right — which is exactly the property one wants in the thing that took the most to build.

Second, the **methodology and the engine are held as distinct intellectual property**. FRAME™ is a trademark of the Healthcare Rebel Alliance; Kincaid IQ is an asset of SiriusB. The methodology that decides what a fiduciary must do is legally separable from the engine that supplies the evidence — which is the precise mirror, in trademark and corporate form, of the substrate-versus-methodology distinction drawn in Section V. The separation of layers, the separation of entities, and the separation of liability are three views of one decision.

The Synovient relationship is included here for completeness and bounded carefully, in keeping with the discipline of naming scope precisely. It is a *Phase-2 architectural extension* for Fiduciary Insights only, not part of the present MVP, and it carries two distinct tokenization roles: titling plan data under CAA §724, and protecting the FRAME™ methodology as intellectual property. The governing instruction is exact — the partner must not preclude the extension, but must not build it ahead of its time. It is named so that the architecture leaves room for it; it is deferred so that the room stays empty until the foundation beneath it is stable.

Why the separation is *the point*.

The two-layer architecture is, finally, a moral architecture — and the person it protects is the one who appears in neither product.

It would be possible to read everything above as an exercise in engineering taste — a tidy separation of concerns, satisfying to architects, of no particular consequence to anyone else. That reading would miss the point entirely. The separation of engine and instrument is, in the end, a *moral* architecture, and the clearest way to see this is to ask who the system is actually for.

It is not, finally, for the committee chair, or the broker, or the consultant, though each of them touches it. It is for the plan participant — the worker and the dependent whose names appear nowhere on a benefits-committee agenda, who will never log into Fiduciary Insights or call the Kincaid IQ API, and whose paycheck deduction is nonetheless the entire reason either system exists. Roughly eighty million Americans pay too much for too few drugs because the intermediaries standing between them and the manufacturers have, for two decades, been quietly smarter than the people charged with overseeing them. The plan sponsors charged with that oversight have wanted to do better and have lacked the analytical machinery to do it. Their participants have borne the cost, in dollars they never saw itemized and choices they never knew were made on their behalf.

This is the register in which the architecture's central artifact must be understood. The **running fiduciary record is not an audit log**. It is a participant-protection artifact. An audit log exists to satisfy a process; the running record exists so that a participant is protected by something more durable than a promise — by a defensible account of every decision made in their name, reproducible years later, when it matters most. That framing is not incidental to the product design; it shapes it. It is why the audit ledger writes from the first day, why every analysis carries provenance through to the engine, and why a 2026 determination is built to survive a 2031 subpoena.

Seen this way, each property of the two-layer system is a way of keeping faith with a person who is not in the room. The engine's litigation-grade provenance means a fiduciary decision can actually be *defended* — so the participant is protected by a record rather than a reassurance. The instrument's dual-buyer alignment means the tool resists being turned against the participant even by the broker who sells it — so the participant is protected by structure rather than by goodwill. The quarterly cadence means oversight is continuous rather than annual — so the participant is watched over between meetings, in the long stretches when no one is looking. None of these protections would hold as well in a single collapsed system, where the deep rigor the engine owes a magistrate and the human immediacy the instrument owes a committee would compromise each other. The separation is what lets the machinery be rigorous and humane at once.

That is the deepest answer to the question this paper set out to address. Kincaid IQ and Rebel IQ are not two products that do similar things. They are two halves of a single instrument of stewardship — one that reasons with the cold precision a courtroom will someday demand, and one that delivers that reasoning with the care a human decision requires — built, in the end, for someone who will never see either, and who is the entire point of both.

Synthesis — a comparison, and a *conclusion*.

The distinctions, drawn together; and a single sentence to carry away.

The differences accumulated across this paper resolve into a clean dichotomy, set out below. The table is a summary, not a substitute — the reasons behind each row are the substance, and they are above — but it is useful to see the whole shape at once.

DIMENSION	KINCAID IQ — THE ENGINE	REBEL IQ — THE INSTRUMENT
ROLE	Intelligence substrate. Reasons about fiduciary state.	Application & delivery. Renders reasoning into a decision and a record.
OWNER	SiriusB AI Data Sciences (IP-holding parent).	Healthcare Rebel Alliance (operating brand).
INTERFACE	Versioned, authenticated API. Consumed as a service.	Committee platform + per-engagement consulting documents.
AGENTS	Ten substrate agents (PCIA, FRSA, FLQA, RADA, CCDA, SEMA, HLSA, ENGA, BOIA, GAA).	FRAME™ six-agent methodology, orchestrating the engine through three gates.
CADENCE	Continuous — every minute.	Quarterly (Fiduciary Insights); per-engagement (the two wedges).
BUYER	Substrate-neutral — whoever holds the obligation.	Two-sided — brokers (offensive) and plan sponsors (defensive).
BUILDS	The reasoning: extraction, scoring, quantification, detection.	The experience, the fiduciary memory, the disclosure, the record.
LIABILITY	Deep intelligence surface — its own to defend.	Isolated application surface — provenance carried from the engine.
THE PARTICIPANT	Protected by provenance — a defensible record.	Protected by alignment — a tool that resists misuse.

What the table cannot show is the most important thing, which is that the two columns are not alternatives. They are not two answers to one question; they are two answers to two questions — *what is true about this plan's fiduciary state*, and *what must this fiduciary do about it* — that a single system would have answered worse by trying to answer together. The engine is unremarkable in isolation, as any reasoning substrate is; the instrument is unremarkable in isolation, as any committee tool is. What is not unremarkable is the combination, executed with the discipline to keep the seam between them clean.

Kincaid IQ reasons. Rebel IQ delivers. Neither is the other, and the refusal to collapse them is not a compromise but the design — the one decision from which the rigor, the alignment, the defensibility, and finally the humanity of the whole system descend. The engine and the instrument: built apart, on purpose, so that together they can keep faith with the person who is the point of both.

To work.

Notes & authorities.

The governing regime against which this architecture is located — cited to fix the system's claims in law, not to advise.

ERISA — THE FIDUCIARY STANDARD

ERISA §404(a)	Prudence and loyalty. The standard of conduct the entire system exists to serve — every gate, score, and record resolves back to it.
ERISA §408(b)(2)	Reasonable compensation for services. The statutory substance of the Reasonableness gate (G2) and of the Cost Plus unit comparison.
ERISA §406	Prohibited transactions and party-in-interest dealing. The substance of the Prohibited-Transaction gate (G3) and its three-layer test.
ERISA §3(21)	Functional fiduciary status — who is bound by the duty, determined by conduct rather than title.

CONSOLIDATED APPROPRIATIONS ACT, 2021 — TRANSPARENCY & DISCLOSURE

CAA §408(b)(2)(B)	Broker and consultant compensation disclosure. The precondition of platform access on which the dual-buyer alignment depends.
CAA §202 · §204	Pharmacy-benefit, drug-cost, and prescription-drug reporting — the disclosure surface the Service-Provider agent reads against.
CAA §724	Plan-data title provisions. The statutory basis for the Phase-2 tokenization architecture; outside MVP scope, preserved by design.

CONTROLLING & PERSUASIVE AUTHORITY

Cunningham v. Cornell (U.S. 2025)	The pleading standard for §406 prohibited-transaction claims — why existence findings are kept separate from content findings.
Tiara Yachts (6th Cir. 2025)	Fiduciary status and the handling of plan assets — a marker for where the obligation attaches.
EBSA FAB 2026-01	Field Assistance Bulletin establishing current enforcement posture — the climate the monitoring infrastructure anticipates.
Arnold et al., Health Affairs (Nov. 2025)	Empirical evidence on intermediary cost and pharmacy-benefit economics — the ground truth beneath the leakage quantification.

This is a positioning and architecture paper, not a legal opinion. Authorities are cited to locate the system against the regime it serves — not to advise any plan, sponsor, or fiduciary, and nothing here substitutes for counsel. The participant whose protection is the object of the whole apparatus appears in none of these citations by name; that absence is the reason the work exists.

CANON NOTE

This paper is written to current canon, in which **Rebel IQ™** denotes the consumer-facing product *family* under the Healthcare Rebel Alliance — with **Rebel IQ™ Fiduciary Insights** as the flagship FRAME™ platform — and **Kincaid IQ** denotes the upstream engine held by SiriusB AI Data Sciences. The retired single-product name “Rebel Fee IQ” still appears in earlier build documents (the **state-of-the-build** memorandum); its scope is now carried by the family above, principally by Fee Intelligence and Fiduciary Insights. One naming tension is flagged rather than silently resolved: where legacy material treats “Rebel Fee IQ” as a standalone fee product, this paper treats fee analysis as one wedge within a broader family. Ratify or correct at the principal's discretion.